



Data Persistence

MICROSERVICES



Content

- Overview
- Pattern and Anti-Pattern
- Caching
- Event Sourcing / CQRS / Sagas

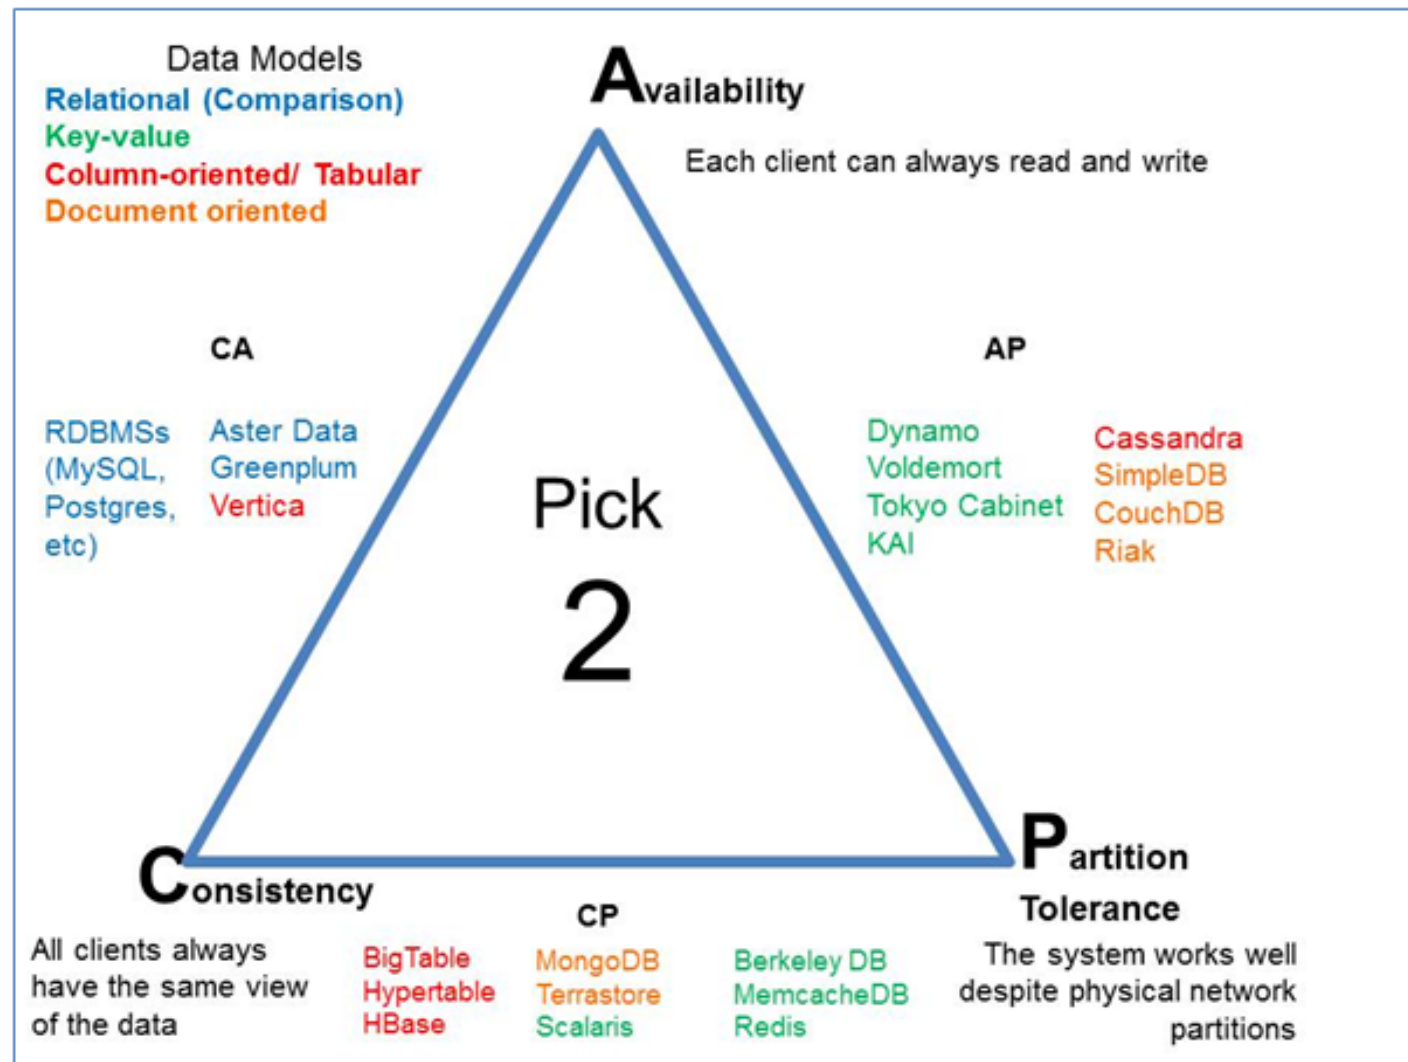


Overview

- People try to copy Netflix, but they can only copy what they see. They copy the results, not the process. — Adrian Cockcroft, former Chief Cloud Architect, Netflix
- Data is the hardest part in microservices!
- CRUD (Create, Read, Update, Delete) is not enough for microservices
- You cannot do ACID (atomicity, consistency, isolation, durability) over multiple datasources transactions → Better use BASE (Basically Available, Soft state, Eventual consistency)



CAP




➔ More @Datenbanken 2 (Master)

<http://nosql-guide.com/wp-content/uploads/1970/01/CAP-Theorem.png>



Pattern and Anti-Pattern


- Anti-Pattern: Shared Database
 - While microservices appear independent, transitive dependencies in the data tier all but eliminate their autonomy





Pattern and Anti-Pattern


- Pattern: Data API
 - Microservices do not access data layer directly
 - Expect for the microservice that implement the data API
 - A Surface area to
 - Implement access control
 - Implementing throttling
 - Perform logging
 - Other policies
 - Possibly coupled with Parallel deployment





Pattern and Anti-Pattern

- Pattern: Bounded Context
 - Domain-Driven-Design
 - Each bounded context has a single, unified model
 - Relationships between models are explicitly defined
 - A product team usually has a strong correlation to a bounded context
 - Ideal pattern for Data APIs – do not fall into the trap of simply projecting current data models
 - Model transactional boundaries as aggregates
 - ➔ Focus on domain models, not data models




More in presentation from 07.12



Pattern and Anti-Pattern

- Pattern: Database per Service and Client Side Joins
 - Support polyglot persistence
 - Independent availability, backup/restore, access patterns, etc.
 - Services must be loosely coupled so that they can be developed, deployed and scaled independently

➔ Microservices need a Cache! & Materialized Views



https://www.slideshare.net/ceposta/the-hardest-part-of-microservices-your-data?qid=249549ff-4a01-4951-8565-9fbffceb7805&v=&b=&from_search=2




Caching

■ Caching Patterns


■ Look Aside

- Attempt retrieval from cache
- Client retrieves from source
- Write into cache



■ Read-through

- Attempt retrieval from cache
- Cache retrieves from source and stores in cache
- Return value to client






Caching

■ Caching Patterns


■ Write-through

- Write to cache
- Cache writes to source
- Ack sent to client



■ Write-behind

- Write to cache
- Ack sent to client
- Cache writes to source asynchronously





Caching


- Requirements
 - Distributed
 - Over various failure boundaries AZ (Availability Zones), Regions (Data Centers)
 - Data replication
 - Tunable consistency
 - Available
 - Multi-node
 - Recovery process protects against any data loss
 - Scalable
 - In-memory performance
 - Horizontally scalable
 - Ease of Provisioning

- ➔ So perhaps think about the pattern: Cache Warming



Two Generals Problem

- Sooner or later you will face the two generals problem in a microservice architecture
- So you will see that CRUD is the wrong approach for microservices
- It is proven to be unsolvable, so you need another approach
- In different databases the application cannot simply use a local ACID transaction



<https://www.youtube.com/watch?v=holjbuSbv3k>



Event Sourcing

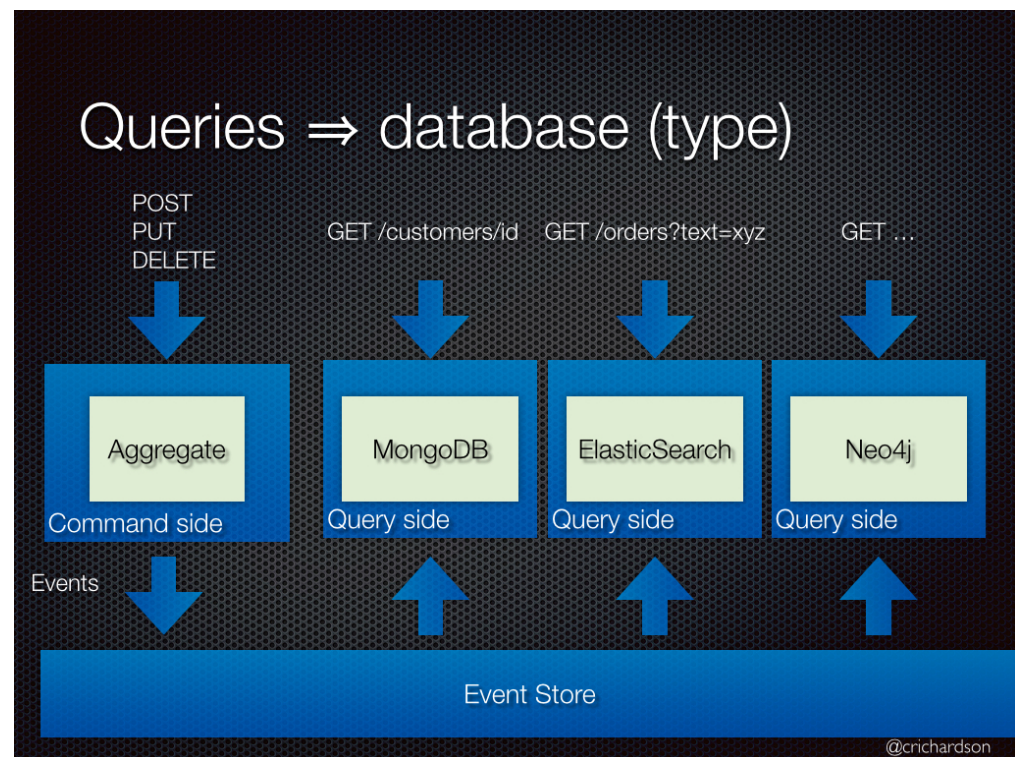
- Event sourcing persists the state of a business entity such as an Order or a Customer as a sequence of state-changing events. Whenever the state of a business entity changes, a new event is appended to the list of events. Since saving an event is a single operation, it is inherently atomic. The application reconstructs an entity's current state by replaying the events.
- Applications persist events in an event store, which is a database of events. The store has an API for adding and retrieving an entity's events. The event store also behaves like a message broker. It provides an API that enables services to subscribe to events. When a service saves an event in the event store, it is delivered to all interested subscribers.
- Some entities, such as a Customer, can have a large number of events. In order to optimize loading, an application can periodically save a snapshot of an entity's current state. To reconstruct the current state, the application finds the most recent snapshot and the events that have occurred since that snapshot. As a result, there are fewer events to replay.
- ➔ Good Frameworks for doing this: akka persistence or lagom

<http://microservices.io/patterns/data/event-sourcing.html>



Command Query Responsibility Segregation (CQRS)

- Split the application into two parts: the command-side and the query-side. The command-side handles create, update, and delete requests and emits events when data changes. The query-side handles queries by executing them against one or more materialized views that are kept up to date by subscribing to the stream of events emitted when data changes.

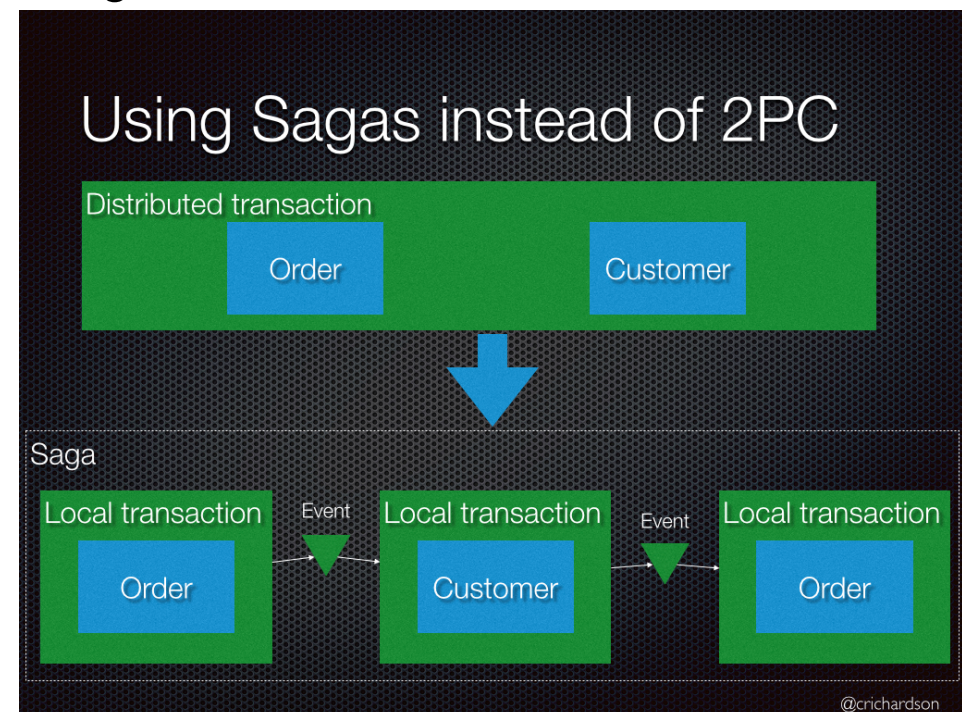


<http://microservices.io/patterns/data/cqrs.html>



Saga

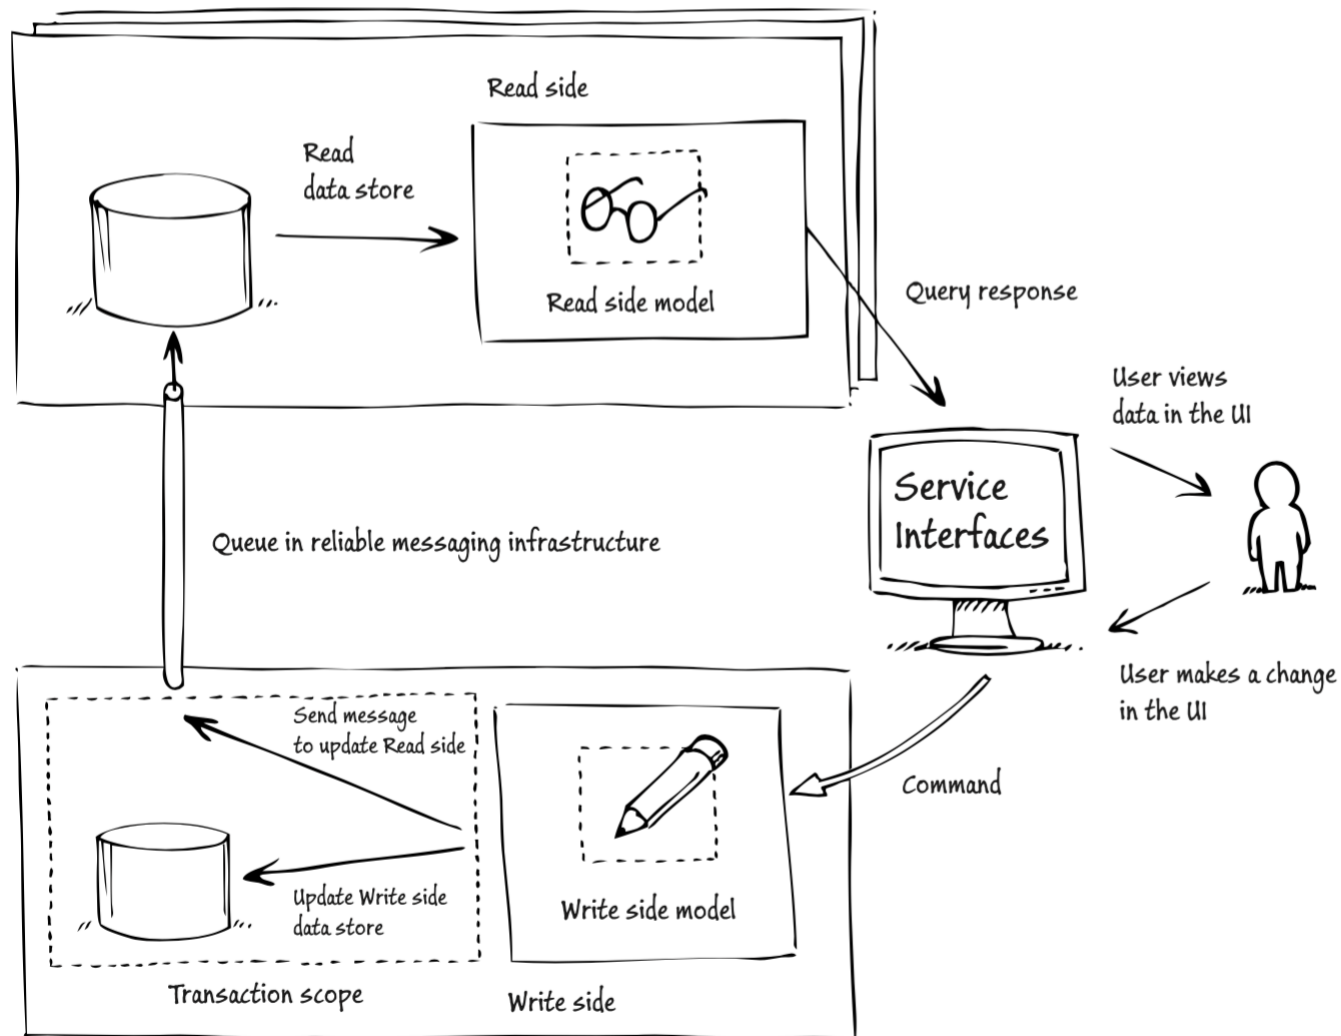
- Implement each business transaction that spans multiple services as a saga. A saga is a sequence of local transactions. Each local transaction updates the database and publishes a message or event to trigger the next local transaction in the saga. If a local transaction fails because it violates a business rule then the saga executes a series of compensating transactions that undo the changes that were made by the preceding local transactions.



<http://microservices.io/patterns/data/saga.html>



Event Sourcing and CQRS



<https://msdn.microsoft.com/en-us/library/jj554200.aspx>

<https://www.slideshare.net/jboner/the-road-to-akka-cluster-and-beyond>